

# Design Document: multi-threaded rpcserver

Perry David Ralston Jr.

CruzID: pdralsto

November 18, 2020

## 1 Goals

This multi-threaded RPC Server will build off of the design of the previous RPC server. The design for the original RPC Server can be found at the end of this document. The design from there remains unchanged unless specifically called out in this design document. The goal of this project is to create an RPC server capable of handling simultaneous requests from numerous clients, leveraging the synchronization techniques learned in lecture. Noted additions to this server is the dispatch thread, responsible for assigning worker threads to the inbound connections, and a shared key-value store for managing the store of variable values that all of the threads can access and modify in a semi-volatile manner. The key-value store persists across all connections to the server and is readable/editable by any thread with coordination.


## 2 Hash Table

For this project, I will be reviving and heavily modifying a hash table that I coded in March of 2017. The key structure of the hash table is a fixed size array of Node\*(s) that make up a linked list. The nodes themselves contain a name stored as a char\* and a value, int64\_t. The hashing is done using a basic hashing function with a fixed decimal constant. The hash table supports the following public functions: insert, replace, delete, clear, dump, and load. The algorithms are detailed below.

### 2.1 Node Structure and Functions

Nodes are structured as a key value pair with a pointer to the next node in the list. Nodes are inserted into the hash table by hashing their key to

acquire the index for the list that they will be inserted into. See diagram below.



**Singly Linked List**

//TODO figure out the synchronization process.

## 2.2 Hash Table Functions

```
1 genHash
2 Input uint8_t* key: string key to be hashed
3 Return int32_t: hashed_value corresponding this key's linked
   list location
4   iterate string and add all char values together
5   store sum in key
6   return floor(tblSize * ((key * HASHCONST) - floor(key *
   HASHCONST)))
```

**Hashing Function**

```
1 Insert
2 Input uint8_t* key: key to insert into this hash table
3 Input int64_t value: value assigned to key
4     hash = genHash(key)
5     make new_node containing key and value
6     if this[hash] is empty
7         this[hash] = new_node
8     else
9         parse this[hash]
10            if next is null
11                next = new_node
12                return
13            if next.key == key
14                temp = next
15                next = new_node
16                new_node.next = temp.next
17                free(temp)
18            return
```

#### Insert

```
1 replace
2 Input uint8_t* key: key to replace in this hash table
3 Input int64_t value: value assigned to key
4     insert(key, size, value)
```

#### Replace

```
1 remove
2 Input uint8_t* key: key to delete from the table
3     hash = genHash(key)
4     if hash_table[hash] is empty
5         return
6     current = hash_table[hash]
7     while current != null
8         if current.key == key
9             if follower == null
10                 hash_table[hash] = current.next
11             else
12                 follower.next = current.next
13             delete current
14         return
```

#### Remove

```
1 lookup
2 Input uint8_t* key: key to find a value for
3 Input int64_t* value: Where to store the found value
4 Return int8_t: 0 if value was found successfully, -1 otherwise
5     hash = gen_hash(key)
6     parse this[hash]
7         if current.key == key
8             value = current.value
9             return 0
10    return -1
```

### Lookup

```
1 clear
2     parse hash_table
3         current = hash_table[i]
4         while current != null
5             leader = current.next
6             delete current
7             current = leader
8     hash_table[i] = null
```

### Clear

```
1 dump
2 Input const char* filename: the file to save the contents of
   this hash_table to
3 Output int8_t: 0 if dump is successful, -1 otherwise, errno is
   set accordingly
4     open(filename)
5     //check for errors
6     parse hash_table
7         current = hash_table[i]
8         while current != null
9             //write key=value to filename
10            //check for errors
11            current = current.next
```


### Dump

```
1 load
2 Input uint8_t* filename: the file to load the contents of this
   hash_table from
3 Output int8_t: 0 if load is successful, -1 otherwise, errno is
   set accordingly
4 open(filename)
5 //check for errors
6 until eof is reached
7     read file into buffer
8     //check for errors
9     while buffer has data
10         extract whitespace seperated value from buffer
11         seperate value on '=' char
12         this.insert(lh_value, rh_value)
```

Load

### 3 Multi-Threading

In the previous rpcserver, request handling was limited to a single request at a time. Using the shared hash table described above and the pthread library, this rpcserver will have the ability to service -N clients at the same time. -N is a command line argument that denotes the number of threads that the server should initially service. The default N value is 4. To achieve this, a structure will be created to house the thread references and act as a means of communication between the dispatch thread and the worker threads. The over all design is pictured below:



## Main Thread {

Create threads[n] where n = -N argv

Create mainMutex, init to running state

Repeat n times

Init threads[i] with a new struct thread

Call pthread\_create, pass it start and threads[i]

Initialize the network

Infinite loop

cl = accept(sock)

While (i = findWorker() == -1) wait(mainMutex)

threads[i].cl = cl

signal(threads[i].mutex)

findWorker  
returns available  
thread or -1

}

## start(void\* arg) {

self = (struct thread\*) arg

While (self->cl == 0) wait(self->mutex)

process(self->CL) //until eof

self->cl = 0

signal(mainMutex)

process is the  
original rpcserver  
code after init

}

This diagram was made using information  
made available in CSE-130 discussion by James Hughes 11/16/2020

# Design Document: rpcserver

Perry David Ralston Jr.

CruzID: pdralsto

October 30, 2020

## 1 Goals

This RPC Server will handle requests from a client to perform simple arithmetic: addition, subtraction, and multiplication (64 bit width) and basic file operations: read, write, create, and filesize query. All messages to the server are assumed to be in network byte order and all responses from the server will be formatted the same.

## 2 Initialization

Server initialization is handled at the command line by running `rpcserver` with the arg `<host_name>:<port>`. Using code retrieved from [CSE130's Canvas Page](#). The command line arg is split at the `:` and the hostname and port are placed in the appropriate positions. The hostname must be  $\leq 1$  kB in size and the server will crash out otherwise. Once the socket has been created and bound, `rpcserver` runs an infinite loop listening on the address and port specified at the command line.

### 2.1 Handling Requests and Sending Responses

Requests from the client are assumed to be in network byte order. Responses will be formatted in network byte order prior to sending responses "over-the-wire". The subroutines shown below detail the process for converting to and from network byte order. This server makes no assumptions about the rate at which all of the data is received from the client. Messages from the client are read into a `bounded_buffer` object for processing.

```

1 Input Bounded_Buffer& bound_buff: reference to a
   Bounded_Buffer object
2 Input var_int_type to_convert: variable width integer type,
   specified at time of call
3 Input int cl: client file descriptor needed for pushByte call
4
5 for i in (sizeof(to_convert) ... 0]
6     to_push = (to_convert >> 8i) & 0xFF
7     bound_buff.pushByte(cl, to_push)

```

#### Convert to Network Byte Order

```

1 Input Bounded_Buffer& nbo_array: Bytes to be converted
2 Input int cl: client file descriptor needed for getByte call
3 Output var_int_type: variable width integer type
4
5 for i in [0 ... sizeof(var_int_type))
6     converted = (converted << 8) + b_buff.getByte(cl)
7 return converted

```

#### Convert from Network Byte Order

## 2.2 Bounded Buffer Class

The Bounded\_Buffer class is responsible for maintaining, filling, and flushing the bounded buffer used to store messages to and from the client for internal processing. It has private members for three uint8\_t pointers that maintain the root, start and end of the buffer and public functions empty, fill, flush, getByte, getBytes, pushByte, and pushBytes. The public functions are detailed below:

```

1 empty()
2 Output boolean: true if there is no data to read from the
3     buffer
4     return start == end
5
6 fill()
7 Input cl: client file descriptor
8 Output: returns bytes read
9     if (start != root)
10         start = root
11     recv up to BUFFER_SIZE number of bytes into buffer from cl
12     if bytes_read == -1
13         log the error and exit
14     end = start + bytes_read + 1
15
16 flush()
17 Input cl: client file descriptor

```



```
18 Output: returns bytes sent
19     if start == end
20         return 0
21     write end - start bytes to cl
22     start = end = root
23     if bytes_sent = -1
24         log the error and exit
25     return bytes_sent
26
27 getByte()
28 Input cl: client file descriptor
29 Output: return pointer to the next byte available in the array
30         or NULL if no bytes are available after attempting
31         to fill the buffer from cl
32     if empty
33         fill
34     if empty
35         return NULL
36     ret_value = start[0]
37     ++start
38     return ret_value
39
40 getBytes()
41 Input cl: client file descriptor
42 Input size_t size: number of bytes to read from the buffer
43 Input uint8_t* dest: destination array
44 Output int8_t: 0 if all requested bytes could be read -1
45                otherwise
46     for i in [0 ... size)
47         currByte = getByte()
48         if currByte == NULL
49             return -1
50         dest[i] = getByte()
51     return 0
52
53 pushByte()
54 Input cl: client file descriptor
55 Input uint8_t in_byte: the byte to be placed in the buffer
56 Output int8_t: 0 if byte was written successfully, -1 otherwise
57     if end > root + BUFFER_SIZE
58         flush(cl)
59     end[0] = in_byte
60     ++end
61
62 pushBytes()
63 Input cl: client file descriptor
64 Input uint8_t* in_bytes: the bytes to be placed in the buffer
65 Input size_t size: the number of bytes to be placed in the
66     BufferError
```

```
65 Output int8_t: 0 if size bytes was written successfully, -1
   otherwise
66     if size < remaining capacity
67         flush
68         if not empty
69             return -1
70     for i in [0 ... size)
71         if pushByte(cl, in_bytes[i]) == -1
72             return -1
73     return 0
```

#### Bounded\_Buffer Public Functions

### 2.3 Resolving Arguments and Calling Functions

Once the request has been parsed, the corresponding function call is made. If no corresponding function call can be found then response header containing **EBADRQC** is sent back to the client. Arguments to the corresponding function are parsed from the data portion of the request.

## 2.4 Supported Functions

**Math Functions** Add, Subtract, and Multiply are supported

**add** Add two numbers, A and B, together returning the value. If overflow would occur set err\_code to EINVAL(22)

```
1 add
2 Input int64_t a: number to add to b
3 Input int64_t b: number to add to a
4 output int64_t: result of a + b
5     set errno to 0
6     if result will overflow
7         set errno to EINVAL
8         return EINVAL
9     result = a + b
10    return result
```

**subtract** Subtract B from A, returning the value. If overflow would occur set err\_code to EINVAL(22)

```
1 subtract
2 Input int64_t a: number to subtract b from
3 Input int64_t b: number to subtract from a
4 output int64_t: result of a - b
5     set errno to 0
6     if result will overflow
7         set errno to EINVAL
8         return EINVAL
9     result = a - b
10    return result
```

**multiply** Add two numbers, A and B, together returning the value. If overflow would occur set err\_code to EINVAL(22)

```
1 multiply
2 Input int64_t a: number to multiply by b
3 Input int64_t b: number to multiply by a
4 output int64_t: result of a * b
5     set errno to 0
6     if result will overflow
7         set errno to EINVAL
8         return EINVAL
9     result = a * b
10    return result
```

**File Functions** Read, Write, Create, and Filesize are supported

**read** Read bufsize bytes from file into buffer starting at the offset and return the number of bytes read if there was no error, -1 otherwise.

```
1 read_file
2 Input char* filename: file to return the size of
3 Input uint64_t offset: where to start reading from
4 Input uint16_t bufsize: how many bytes to read from the file
5
6     file_size = filesize(filename)
7     if file_size == -1
8         set header error status
9         send header
10        return -1
11    if filesize - offset < bufsize
12        set header status
13        send header
14        return -1
15    file_d = open(filename, O_RDONLY)
16    if file_d == -1
17        set header error status
18        send header
19        return -1
20    send header
21    while bytes_read < bufsize
22        curr_read = read(file_d, buffer, BUFFER_SIZE)
23        if curr_read == -1 or == 0
24            close connection
25            return -1
26        if bytes_read + curr_read <= bufsize
27            send curr_read bytes to client
28            bytes_read += curr_read
29        else
30            send bufsize - bytes_read to client
31            return bufsize
32    return 0
```

**write** Write bufsize bytes from buffer to a file starting at offset and return the number of bytes written if there was no error, -1 otherwise.

```
1 write_file
2 Input char* filename: file to return the size of
3 Input uint64_t offset: where to start reading from
4 Input uint16_t bufsize: how many bytes to write to the file
5     file_d = open(filename, O_WRONLY)
6     if file_d == -1
7         set header error status
8         send header
9         return -1
10    send header
11    while bytes_written < bufsize
12        fill buffer up to bufsize - bytes_written
13        if no bytes written to buffer
14            close connection
15            return -1
16        write filled_bytes to filename
17        bytes_written += filled_bytes
18    return 0
```

**create** Create a new 0 byte file if it does not already exist, returns -1 if an error occurs

```
1 create_file
2 Input char* filename: file to create
3 Output int64_t: 0 if successful, -1 otherwise
4     if open(filename, O_CREATE, O_EXCL, O_WRONLY) == -1
5         set header error status
6         send header
7         return -1
8     send header
9     return 0
```

**filesize** Returns the size of an existing file, -1 in the event of an error

```
1 get_filesize
2 Input char* filename: file to return the size of
3     if stat(filename, fileStats) == -1
4         set header error status
5         send header
6         return
7     return fileStats.st_size
```