

Design Document

Persistant Multi-Threaded RPCServer

Perry David Ralston Jr.
CruzID: pdralsto

December 11, 2020

1 Goals

1. Fix the bugs from assignment 2
2. Support synchronized storing of values or variable names in key/value store
3. Support recursive name resolution to a specified max depth
4. Support large scalability
5. Key-Value store will be persistant across instantiations of the server

2 Synchronized Hashtable

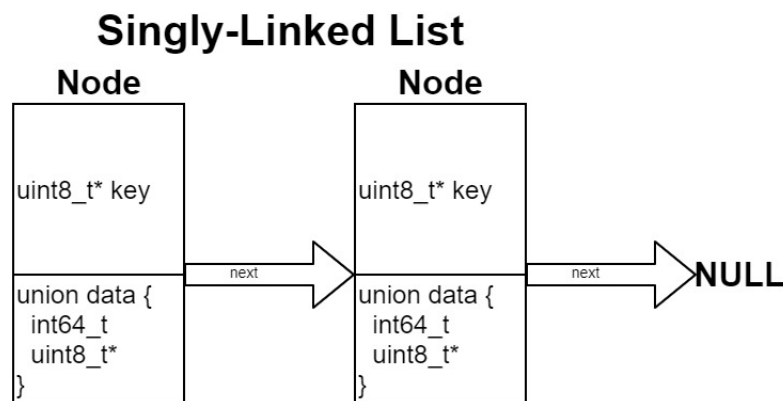
The hashtable defined below utilizes the djb2 hashing function, sourced from <http://www.cse.yorku.ca/~oz/hash.html> on 11/23/2020. It supports these public functions:

insert	remove	lookup
clear	dump	load
acquire	release	


Data is stored as a Singly-Linked List of Node objects, detailed below. The hash table is wrapped by an inherited class, SyncHash, that adds index level synchronization to the table. The synchronization is managed using a parallel array of mutexes and all of the functions use calls to acquire and release to change the states of them.

2.1 Node Structure and Functions

Nodes are structured as a key value pair with a pointer to the next node in the list. Nodes are inserted into the hash table by hashing their key to acquire the index for the list that they will be inserted into. Nodes are instantiated with their next pointers set to null. The data of the node is a union type `int64_t` and `char*`. This allows for the nodes to store either a variable name or a numeric value, supporting recursive name resolution. See diagram below.



Hash-Table of Node*



2.2 Hash Table Functions

```
1 genHash
2 Input uint8_t* key: string key to be hashed
3 //sourced from http://www.cse.yorku.ca/~oz/hash.html 11/23/2020
4 Return int32_t: hashed_value corresponding this keys linked
   list location or -1 if the key is invalid
5   if validate_key(key) == -1
6       return -1
7   key_value = 5381
8   traverse key
9       key_value = key_value * 33 + key[i]
10  return floor(tblSize * ((key * HASHCONST) - floor(key *
    HASHCONST)))
```

Hashing Function

Acquire and release supports two different public interfaces, single and list. The list version acquires locks in index order to avoid dead lock.

```
1 acquire
2 Input uint8_t* key: key to acquire lock on
3 Input int64_t ident: unique thread identifier
4 Returns int8_t: 0 if successful. -1 If the key is invalid
5     hash = genHash(key)
6     if hash == -1
7         return -1
8     acquire(hash, ident)
9     return 0
```

Acquire

```
1 acquire
2 Input uint8_t** keys: keys to acquire locks on
3 Input int64_t ident: unique thread identifier
4 Returns int8_t: 0 if successful. -1 If a key is invalid
5     hashlist = {genHash(key) for key in keys}
6     sort(hashlist)
7     if hashlist[0] == -1
8         return -1
9     for hash in hashlist
10         acquire(hash, ident)
11     return 0
```

Acquire (list)

```
1 release
2 Input uint8_t* key: key to release lock on
3 Input int64_t ident: unique thread identifier
4 Returns int8_t: 0 if successful. -1 If the key is invalid
5     hash = genHash(key)
6     if hash == -1
7         return -1
8     if this thread owns the lock
9         release(hash)
10    return 0
```

Release

```
1 release
2 Input uint8_t** keylist: keys to release lock on
3 Input int64_t ident: unique thread identifier
4 Returns int8_t: 0 if successful. -1 If the key is invalid
```

```
5  hashlist = {genHash(key) for key in keys}
6  if hash == -1
7      return -1
8  if this thread owns the lock
9      release(hash)
10 return 0
11
12 hashlist = {genHash(key) for key in keys}
13 sort(hashlist)
14 if hashlist[0] == -1
15     return -1
16 for hash in hashlist
17     acquire(hash, ident)
18 return 0
```

Release (List)

```
1  Insert
2  Input uint8_t* key: key to insert into this hash table
3  Input int64_t value: value assigned to key
4  Input int64_t ident: thread specific identifier
5  Return int8_t: 0 if insert is successful, -1 otherwise, errno
   is set appropriately
6  hash = acquire(key, ident);
7  if hash == -1
8      return -1
9  current = hashtable[hash]
10 prev = hashtable[hash]
11 while current != null
12     if current.key == key
13         current.data = value
14         release(hash)
15         return 0
16     prev = current
17     current = current.next
18 create a newNode using key and value
19 if prev == current
20     hashtable[hash] = newNode
21 else
22     prev.next = newNode
23 release(hash)
24 return 0
```

Insert

```
1  remove
2  Input uint8_t* key: key to delete from the table
3  Input int64_t ident: thread specific identifier
```

```
4 Return int8_t: 0 if remove is successful, -1 otherwise, errno
   is set appropriately
5   hash = acquire(key, ident)
6   if hash == -1
7       return -1
8   current = hashtable[hash]
9   if current is not null and current.key == key
10      hashtable[hash] = current.next
11      delete current
12      return 0
13   follower = current
14   while current is not null
15       if current.key == key
16           follower.next = current.next
17           delete current
18           release(hash);
19           return 0
20       follower = current
21       current = current.next
22   release(hash)
23   return -1
```

Remove

```
1 lookup
2 Input uint8_t* key: key to find a value for
3 Input int64_t* value: Where to store the found value
4 Input int64_t ident: thread specific identifier
5 Return int8_t: 0 if value was found successfully, -1 otherwise
6   hash = acquire(key, ident)
7   traverse hashtable[hash]
8       if current.key == key
9           value = current.value
10          signal(hash)
11          return 0
12   release(hash)
13   return -1
```

Lookup

```
1 clear
2   acquire all of the locks
3   traverse hashtable
4       if hashtable[i] is defined
5           current = hashtable[i]
6           while current is defined
7               temp = current
8               current = current.next
```

```

9         delete temp
10        hashtable[i] = null
11        release(i)

```

Clear

```

1 dump
2 Input char* filename: file to print the key-value store to
3   acquire all of the locks
4   traverse hashtable
5       if hashtable[i] is defined
6           current = hashtable[i]
7       while current is defined
8           print key=value\n
9   release(i)

```

Dump

```

1 load
2 Input uint8_t* filename: the file to load the contents of this
   hash_table from
3 Input int64_t ident: unique thread identifier
4 Output int8_t: 0 if load is successful, -1 otherwise, errno is
   set accordingly
5   open the file named filename
6   //error handling
7   acquire all locks
8   until eof is reached
9       read a line from the file
10      //error handling
11      parse the line into key and value
12      //error handling
13      insert(key, value, ident)
14  release all locks
15  return 0


```

Load

3 Multi-Threading

In the previous rpcserver, request handling was limited to a single request at a time. Using the shared hash table described above and the pthread library, this rpcserver will have the ability to service -N clients at the same time. -N is a command line argument that denotes the number of threads that the server should initially service. The default N value is 4. To achieve this, a

structure will be created to house the thread references and act as a means of communication between the dispatch thread and the worker threads. The overall design is pictured below:



Threads

The instructions for request handling are moved into a new function, `process`, that receives a `Thread&` as an argument. The `Thread` type is shown above and the members are used by the `process` function in order to facilitate synchronization of the access to the hashtable as well as retrieving and responding to the request from the client.

```

1 process
2 Input Thread* self:  contains all of the information associated
   with this Thread
3 Reliant data:
4   math_funcs:        Array of math function pointers to reduce
   repeated code
5   resp_header:       struct representing the response header
6   buffer:            4kB bounded buffer for client/server data
   handling
7   var_args:          3 index array of uint8_t* to hold keys of
   variable args from client
8   ARG_COUNT:         Constant set to 3 for batch lock
   acquisition of args from the hashtable
9   MATH_OPS:          0x0100
10  FINAL_BYTE:        0X000f
11  MATH_FUNC_COUNT:   5
12  CLEAR_CONFIRM:     0x0badbad0
13  DEL:               0x010f
14  GETV:              0x0108
15  GETV:              0x0109
16  DUMP:              0x0301
17  LOAD:              0x0302
18 Notes: For brevity, error checking is assumed
19   while client connection is open
20   build_header(resp_header, buffer)
  
```



```

21     math_op = (header->op & MATH_OPS) + (header->op &
22     FINAL_BYTE)
23     if math_op != 0
24         fun_index = (math_op & FINAL_BYTE) - 1
25         if fun_index < MATH_FUNC_COUNT
26             Determine the number and which args are variables,
27             assign them to corresponding indicies of var_args
28             if var_args[2] != null
29                 self->htable->acquire(var_args, ARG_COUNT, self->
30                 ident)
31                 if var_arg[0] != null
32                     self->htable->lookup(var_arg[0], arg1, self->ident)
33                 if var_arg[1] != null
34                     self->htable->lookup(var_arg[1], arg2, self->ident)
35                 math_funcs[fun_index](arg1, arg2, result)
36                 if var_args[2] != null
37                     self->htable->insert(var_args[2], result, self->ident
38             )
39             send response
40         else
41             if math_op == DEL
42                 key = read key from buffer
43                 self->htable->remove(key, ident)
44             else if math_op == GETV or SETV
45                 key = read key from buffer
46                 If recursive flag is set
47                     resolved_key = recursive lookup result
48                 else
49                     resolved_key = single stage lookup result
50                 if math_op == GETV
51                     send response
52                 else
53                     key = read key from buffer
54                     if key is valid
55                         self->htable->insert(resolved_key, key, self->
56                     ident)
57                     send response
58                 else
59                     resp_header.error = EINVAL
60                     send response
61             else
62                 resp_header.error = EINVAL
63                 send response
64             if header.op == one of the file operations
65                 behavior is unchanged from asgn1 and 2
66             else if header.op == DUMP
67                 filename = filename retrieved from buffer
68                 self->htable->dump(filename, self->ident)

```

```
65     send response
66     else if header.op == LOAD
67         filename = filename retrieved from buffer
68         self->htable->load(filename, self->ident)
69         send response
70     else if header.op == CLEAR
71         arg1 = pull 32 bit unsigned int from buffer
72         if arg1 == CLEAR_CONFIRM
73             self->htable->clear(self->ident)
74             send response
75         else
76             resp_header.error = EINVAL
77             send response
78     else
79         resp_header.error = EINVAL
80         send response
```

Process

Design Document: rpcserver

Perry David Ralston Jr.

CruzID: pdralsto

October 30, 2020

1 Goals

This RPC Server will handle requests from a client to perform simple arithmetic: addition, subtraction, and multiplication (64 bit width) and basic file operations: read, write, create, and filesize query. All messages to the server are assumed to be in network byte order and all responses from the server will be formatted the same.

2 Initialization

Server initialization is handled at the command line by running `rpcserver` with the arg `<host_name>:<port>`. Using code retrieved from [CSE130's Canvas Page](#). The command line arg is split at the `:` and the hostname and port are placed in the appropriate positions. The hostname must be ≤ 1 kB in size and the server will crash out otherwise. Once the socket has been created and bound, `rpcserver` runs an infinite loop listening on the address and port specified at the command line.

2.1 Handling Requests and Sending Responses

Requests from the client are assumed to be in network byte order. Responses will be formatted in network byte order prior to sending responses "over-the-wire". The subroutines shown below detail the process for converting to and from network byte order. This server makes no assumptions about the rate at which all of the data is received from the client. Messages from the client are read into a `bounded_buffer` object for processing.

```

1 Input Bounded_Buffer& bound_buff: reference to a
   Bounded_Buffer object
2 Input var_int_type to_convert: variable width integer type,
   specified at time of call
3 Input int cl: client file descriptor needed for pushByte call
4
5 for i in (sizeof(to_convert) ... 0]
6     to_push = (to_convert >> 8i) & 0xFF
7     bound_buff.pushByte(cl, to_push)

```

Convert to Network Byte Order

```

1 Input Bounded_Buffer& nbo_array: Bytes to be converted
2 Input int cl: client file descriptor needed for getByte call
3 Output var_int_type: variable width integer type
4
5 for i in [0 ... sizeof(var_int_type))
6     converted = (converted << 8) + b_buff.getByte(cl)
7 return converted

```

Convert from Network Byte Order

2.2 Bounded Buffer Class

The Bounded_Buffer class is responsible for maintaining, filling, and flushing the bounded buffer used to store messages to and from the client for internal processing. It has private members for three uint8_t pointers that maintain the root, start and end of the buffer and public functions empty, fill, flush, getByte, getBytes, pushByte, and pushBytes. The public functions are detailed below:

```

1 empty()
2 Output boolean: true if there is no data to read from the
3     buffer
4     return start == end
5
6 fill()
7 Input cl: client file descriptor
8 Output: returns bytes read
9     if (start != root)
10         start = root
11     recv up to BUFFER_SIZE number of bytes into buffer from cl
12     if bytes_read == -1
13         log the error and exit
14     end = start + bytes_read + 1
15
16 flush()
17 Input cl: client file descriptor

```

```
18 Output: returns bytes sent
19     if start == end
20         return 0
21     write end - start bytes to cl
22     start = end = root
23     if bytes_sent = -1
24         log the error and exit
25     return bytes_sent
26
27 getByte()
28 Input cl: client file descriptor
29 Output: return pointer to the next byte available in the array
30         or NULL if no bytes are available after attempting
31         to fill the buffer from cl
32     if empty
33         fill
34     if empty
35         return NULL
36     ret_value = start[0]
37     ++start
38     return ret_value
39
40 getBytes()
41 Input cl: client file descriptor
42 Input size_t size: number of bytes to read from the buffer
43 Input uint8_t* dest: destination array
44 Output int8_t: 0 if all requested bytes could be read -1
45                otherwise
46     for i in [0 ... size)
47         currByte = getByte()
48         if currByte == NULL
49             return -1
50         dest[i] = getByte()
51     return 0
52
53 pushByte()
54 Input cl: client file descriptor
55 Input uint8_t in_byte: the byte to be placed in the buffer
56 Output int8_t: 0 if byte was written successfully, -1 otherwise
57     if end > root + BUFFER_SIZE
58         flush(cl)
59     end[0] = in_byte
60     ++end
61
62 pushBytes()
63 Input cl: client file descriptor
64 Input uint8_t* in_bytes: the bytes to be placed in the buffer
65 Input size_t size: the number of bytes to be placed in the
66     BufferError
```

```
65 Output int8_t: 0 if size bytes was written successfully, -1
   otherwise
66     if size < remaining capacity
67         flush
68         if not empty
69             return -1
70     for i in [0 ... size)
71         if pushByte(cl, in_bytes[i]) == -1
72             return -1
73     return 0
```

Bounded_Buffer Public Functions

2.3 Resolving Arguments and Calling Functions

Once the request has been parsed, the corresponding function call is made. If no corresponding function call can be found then response header containing **EBADRQC** is sent back to the client. Arguments to the corresponding function are parsed from the data portion of the request.

2.4 Supported Functions

Math Functions Add, Subtract, and Multiply are supported

add Add two numbers, A and B, together returning the value. If overflow would occur set err_code to EINVAL(22)

```
1 add
2 Input int64_t a: number to add to b
3 Input int64_t b: number to add to a
4 output int64_t: result of a + b
5     set errno to 0
6     if result will overflow
7         set errno to EINVAL
8         return EINVAL
9     result = a + b
10    return result
```

subtract Subtract B from A, returning the value. If overflow would occur set err_code to EINVAL(22)

```
1 subtract
2 Input int64_t a: number to subtract b from
3 Input int64_t b: number to subtract from a
4 output int64_t: result of a - b
5     set errno to 0
6     if result will overflow
7         set errno to EINVAL
8         return EINVAL
9     result = a - b
10    return result
```

multiply Add two numbers, A and B, together returning the value. If overflow would occur set err_code to EINVAL(22)

```
1 multiply
2 Input int64_t a: number to multiply by b
3 Input int64_t b: number to multiply by a
4 output int64_t: result of a * b
5     set errno to 0
6     if result will overflow
7         set errno to EINVAL
8         return EINVAL
9     result = a * b
10    return result
```

File Functions Read, Write, Create, and Filesize are supported

read Read bufsize bytes from file into buffer starting at the offset and return the number of bytes read if there was no error, -1 otherwise.

```
1 read_file
2 Input char* filename: file to return the size of
3 Input uint64_t offset: where to start reading from
4 Input uint16_t bufsize: how many bytes to read from the file
5
6     file_size = filesize(filename)
7     if file_size == -1
8         set header error status
9         send header
10        return -1
11    if filesize - offset < bufsize
12        set header status
13        send header
14        return -1
15    file_d = open(filename, O_RDONLY)
16    if file_d == -1
17        set header error status
18        send header
19        return -1
20    send header
21    while bytes_read < bufsize
22        curr_read = read(file_d, buffer, BUFFER_SIZE)
23        if curr_read == -1 or == 0
24            close connection
25            return -1
26        if bytes_read + curr_read <= bufsize
27            send curr_read bytes to client
28            bytes_read += curr_read
29        else
30            send bufsize - bytes_read to client
31            return bufsize
32    return 0
```


write Write bufsize bytes from buffer to a file starting at offset and return the number of bytes written if there was no error, -1 otherwise.

```
1 write_file
2 Input char* filename: file to return the size of
3 Input uint64_t offset: where to start reading from
4 Input uint16_t bufsize: how many bytes to write to the file
5     file_d = open(filename, O_WRONLY)
6     if file_d == -1
7         set header error status
8         send header
9         return -1
10    send header
11    while bytes_written < bufsize
12        fill buffer up to bufsize - bytes_written
13        if no bytes written to buffer
14            close connection
15            return -1
16        write filled_bytes to filename
17        bytes_written += filled_bytes
18    return 0
```

create Create a new 0 byte file if it does not already exist, returns -1 if an error occurs

```
1 create_file
2 Input char* filename: file to create
3 Output int64_t: 0 if successful, -1 otherwise
4     if open(filename, O_CREATE, O_EXCL, O_WRONLY) == -1
5         set header error status
6         send header
7         return -1
8     send header
9     return 0
```

filesize Returns the size of an existing file, -1 in the event of an error

```
1 get_filesize
2 Input char* filename: file to return the size of
3     if stat(filename, fileStats) == -1
4         set header error status
5         send header
6         return
7     return fileStats.st_size
```